

AI 539

Homework 4

Ameyassh Nagarajan

June 2024

1 Section 1

1.1 Task 1.1

To achieve $a \approx v_j$, the dot product $q \cdot k_j$ must be significantly larger than $q \cdot k_i$ for all $i \neq j$. This ensures that the attention weight α_j corresponding to v_j is close to 1, while the other weights are close to 0. Therefore, the query q must be closely aligned with the key k_j .

1.2 Task 1.2

Given that the keys are orthogonal unit vectors, we can construct a query vector q that is equidistant from k_a and k_b . A simple way to achieve this is to take the sum of k_a and k_b :

$$q = k_a + k_b$$

The dot products for this query vector are:

$$q \cdot k_a = (k_a + k_b) \cdot k_a = 1 + 0 = 1$$

$$q \cdot k_b = (k_a + k_b) \cdot k_b = 0 + 1 = 1$$

For any other key k_i orthogonal to k_a and k_b :

$$q \cdot k_i = (k_a + k_b) \cdot k_i = 0 + 0 = 0$$

The softmax function will assign equal weights to k_a and k_b and a much smaller weight to all other keys. Hence, the attention weights will be approximately:

$$\alpha_a \approx \frac{1}{2}, \quad \alpha_b \approx \frac{1}{2}, \quad \alpha_i \approx 0 \quad \text{for all } i \neq a, b$$

Therefore, the output will be:

$$a \approx \alpha_a v_a + \alpha_b v_b = \frac{1}{2}v_a + \frac{1}{2}v_b = \frac{1}{2}(v_a + v_b)$$

1.3 Task 1.3

Given the query vector $q = \mu_a + \mu_b$:

1. For k_a :

$$q \cdot k_a = (\mu_a + \mu_b) \cdot (\mu_a \lambda_a) = \mu_a \cdot \mu_a \lambda_a + \mu_b \cdot \mu_a \lambda_a = \lambda_a$$

This is because $\mu_a \cdot \mu_a = 1$ and $\mu_b \cdot \mu_a = 0$ (due to orthogonality).

2. For k_b :

$$q \cdot k_b = (\mu_a + \mu_b) \cdot (\mu_b \lambda_b) = \mu_a \cdot \mu_b \lambda_b + \mu_b \cdot \mu_b \lambda_b = \lambda_b$$

This is because $\mu_b \cdot \mu_b = 1$ and $\mu_a \cdot \mu_b = 0$ (due to orthogonality).

For any other key k_i orthogonal to μ_a and μ_b :

$$q \cdot k_i = (\mu_a + \mu_b) \cdot (\mu_i \lambda_i) = \mu_a \cdot \mu_i \lambda_i + \mu_b \cdot \mu_i \lambda_i = 0$$

This is because $\mu_a \cdot \mu_i = 0$ and $\mu_b \cdot \mu_i = 0$ (due to orthogonality).

The attention weights are computed using the softmax function:

$$\alpha_i = \frac{\exp(q \cdot k_i / \sqrt{d})}{\sum_{j=1}^m \exp(q \cdot k_j / \sqrt{d})}$$

For k_a and k_b , the dot products are λ_a and λ_b , respectively. Thus, the attention weights α_a and α_b are:

$$\alpha_a \approx \frac{\exp(\lambda_a / \sqrt{d})}{\exp(\lambda_a / \sqrt{d}) + \exp(\lambda_b / \sqrt{d})}$$

$$\alpha_b \approx \frac{\exp(\lambda_b / \sqrt{d})}{\exp(\lambda_a / \sqrt{d}) + \exp(\lambda_b / \sqrt{d})}$$

Since $\lambda_i \sim \mathcal{N}(1, \beta)$, the values of λ_a and λ_b will fluctuate due to the random sampling. This means the attention weights α_a and α_b will also fluctuate, causing variability in the output a .

Over multiple resamplings of $\lambda_1, \dots, \lambda_m$, the output a will vary due to the fluctuations in the attention weights. However, because λ_i values are sampled from a normal distribution centered around 1, the fluctuations will average out over many resamplings. This means that on average, the output a will be approximately:

$$a \approx \frac{1}{2}(v_a + v_b)$$

But individual outputs will show variation due to the noise introduced by the λ 's.

1.4 Task 1.4

In Task 1.2, we designed a query q that would lead to the output being an average of v_a and v_b . For multi-head attention, we need to design two queries q_1 and q_2 that will result in a_1 and a_2 being close to v_a and v_b , and their average should be close to $\frac{1}{2}(v_a + v_b)$.

Given the orthogonal unit vectors μ_i and the noisy scaling factors λ_i , we can construct the queries as follows:

1. First query q_1 should be the sum of μ_a and μ_b :

$$q_1 = \mu_a + \mu_b$$

2. Second query q_2 should be the difference of μ_a and μ_b :

$$q_2 = \mu_a - \mu_b$$

For q_1 :

$$q_1 \cdot k_a = (\mu_a + \mu_b) \cdot (\mu_a \lambda_a) = \lambda_a$$

$$q_1 \cdot k_b = (\mu_a + \mu_b) \cdot (\mu_b \lambda_b) = \lambda_b$$

For q_2 :

$$q_2 \cdot k_a = (\mu_a - \mu_b) \cdot (\mu_a \lambda_a) = \lambda_a$$

$$q_2 \cdot k_b = (\mu_a - \mu_b) \cdot (\mu_b \lambda_b) = -\lambda_b$$

The attention weights for each head will fluctuate due to the noise introduced by the λ 's, but the combination of both heads should average out the noise to some extent.

The outputs for the heads will be:

$$a_1 \approx \frac{1}{2}(v_a + v_b)$$

$$a_2 \approx \frac{1}{2}(v_a - v_b)$$

Combining these:

$$a = \frac{1}{2}(a_1 + a_2) \approx \frac{1}{2} \left(\frac{1}{2}(v_a + v_b) + \frac{1}{2}(v_a - v_b) \right) = \frac{1}{2} \left(v_a + \frac{v_a - v_b}{2} \right) = \frac{1}{2}(v_a + v_b)$$

Thus, the queries q_1 and q_2 designed as $q_1 = \mu_a + \mu_b$ and $q_2 = \mu_a - \mu_b$ will result in the final output being approximately $\frac{1}{2}(v_a + v_b)$.

2 Section 2

2.1 Task 2.1

Here is the implementation of the `SingleQueryScaledDotProductAttention` class in Python:

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4
5 class SingleQueryScaledDotProductAttention(nn.Module):
6     def __init__(self, enc_hid_dim, dec_hid_dim, kq_dim=64):
7         super().__init__()
8         self.q_proj = nn.Linear(dec_hid_dim, kq_dim)
9         self.k_proj = nn.Linear(enc_hid_dim * 2, kq_dim)
10        self.v_proj = nn.Linear(enc_hid_dim * 2, enc_hid_dim
11                                * 2)
12
13        self.scale = torch.sqrt(torch.FloatTensor([kq_dim]))
14        .to(torch.device("cuda" if torch.cuda.
15                        is_available() else "cpu"))
16
17    def forward(self, hidden, encoder_outputs):
18        queries = self.q_proj(hidden).unsqueeze(1) # [
19            batch_size, 1, kq_dim]
20        keys = self.k_proj(encoder_outputs) # [batch_size,
21            src_len, kq_dim]
22        values = self.v_proj(encoder_outputs) # [batch_size
23            , src_len, enc_hid_dim * 2]
24
25        # Compute the dot products
26        energy = torch.bmm(queries, keys.transpose(1, 2)) /
27            self.scale # [batch_size, 1, src_len]
28        alpha = F.softmax(energy, dim=-1) # [batch_size, 1,
29            src_len]
30
31        # Compute the weighted sum of values
32        attended_val = torch.bmm(alpha, values).squeeze(1)
33            # [batch_size, enc_hid_dim * 2]
34
35        return attended_val, alpha.squeeze(1)
```

Listing 1: `SingleQueryScaledDotProductAttention` Class

TASK 2.2 Attention Diagrams

Attention Diagram Examples

In this section, we present a few examples of attention diagrams produced by our model. We will analyze these diagrams to identify common patterns and

discuss their implications. Note that German is (mostly) a Subject-Object-Verb language, so you may find attention patterns that indicate inversion of word order when translating to Subject-Verb-Object English, as seen in the second example below.

Example 1: This diagram shows a clear alignment between the source and target sentences. The attention focuses on specific pairs of words, shows a direct translation.




Figure 1: Example 1

Example 2: The attention shifts, highlighting the Subject-Object-Verb to Subject-Verb-Object transition between German and English.




Figure 2: Example 2

Example 3: This example illustrates how the model handles multiple objects in a sentence, with attention spread across relevant words.

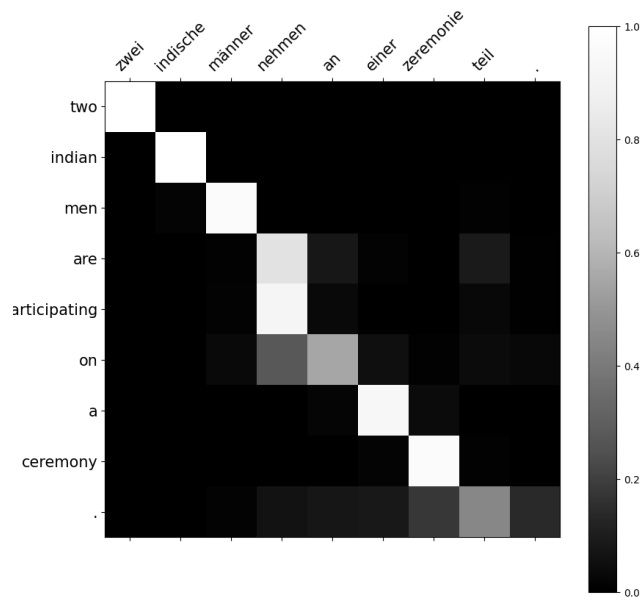


Figure 3: Example 3

Example 4: We observe a strong diagonal pattern, shows a one-to-one correspondence between source and target words.




Figure 4: Example 4

Example 5: The model correctly aligns prepositions and their corresponding objects, showcasing the handling of syntactic structures.




Figure 5: Example 5

Example 6: Attention is spread across several words, showing the model’s capability to handle complex sentence structures.




Figure 6: Example 6

Example 7: Similar to Example 2, this diagram demonstrates the inversion of word order in translation.




Figure 7: Example 7

Example 8: The attention focuses on key nouns and verbs, ensuring accurate translation of the main sentence components.




Figure 8: Example 8

2.2 Task 2.3

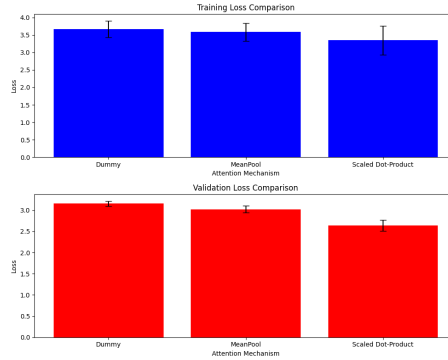


Figure 9: Enter Caption

Training Loss Comparison:

From the training loss plot, we observe the following:

- **Dummy Attention:**
 - Mean Training Loss: Approximately 3.67
 - Variance: 0.232
- **MeanPool Attention:**
 - Mean Training Loss: Approximately 3.58
 - Variance: 0.259
- **Scaled Dot-Product Attention:**
 - Mean Training Loss: Approximately 3.35
 - Variance: 0.411

Validation Loss Comparison:

From the validation loss plot, we observe the following:

- **Dummy Attention:**
 - Mean Validation Loss: Approximately 3.16
 - Variance: 0.059
- **MeanPool Attention:**
 - Mean Validation Loss: Approximately 3.02
 - Variance: 0.081
- **Scaled Dot-Product Attention:**

- Mean Validation Loss: Approximately 2.64
- Variance: 0.129

Observed Trends:

- **Training Loss:**

- The Scaled Dot-Product attention mechanism shows the lowest mean training loss (3.35) compared to Dummy and MeanPool, shows better training performance.
- However, the variance for Scaled Dot-Product is higher (0.411), suggesting more variability in the training process.

- **Validation Loss:**

- The Scaled Dot-Product attention mechanism also shows the lowest mean validation loss (2.64), suggesting better generalization to unseen data.
- Similar to training loss, the variance is higher for Scaled Dot-Product (0.129), shows some variability in the validation performance.

- **Comparison between Dummy and MeanPool:**

- Both Dummy and MeanPool show similar mean training losses (3.67 and 3.58, respectively) with relatively low variance.
- MeanPool slightly outperforms Dummy when it comes to training and validation losses, showing a modest improvement in performance.

- **Overall Performance:**

- The Scaled Dot-Product attention mechanism consistently outperforms Dummy and MeanPool in both training and validation loss metrics.
- The higher variance observed with Scaled Dot-Product shows that while it performs better on average, it may not be stable across different runs.

Conclusion:

The Scaled Dot-Product attention mechanism demonstrates better performance in reducing both training and validation losses compared to Dummy and MeanPool attention mechanisms. However, this results in higher variance, showing that the performance can change more between different runs. MeanPool shows a slight improvement over Dummy, but not as good as Scaled Dot-Product.